

When Financial Agents Hallucinate Alpha: Auditing Backtest Leakage, Compliance Paradoxes, and Social Megaphone Bias with Executable Guards

Howard Ye^{1*} and Shwai He^{2*}

^{1,2}Open Quantitative Integrity Initiative (fin-skills & Multimodal Financial NLP Group)
{howard, shwaihe}@fin-skills.org (*Equal contribution)

Abstract

Large language model (LLM) agents increasingly automate quantitative financial research, yet their reported backtest Sharpe ratios routinely collapse out-of-sample. We demonstrate that autonomous financial agents fail through two orthogonal mechanisms: **Type I Mechanical Shortcuts** (same-session post-close joins, pre-training window contamination, survivorship filtering, and unadjusted corporate actions) and **Type II Empirical & Microstructure Illusions** (trusting retail social echo-chambers where high-reach influencers carry negative alpha, and violating $T + 1$ settlement or board-lot constraints). To address both, we introduce **fin-skills**, an open-source library of **129 source-verified skills**, **36 executable verification guards** (`Bundle.check()`), and **53 MCP agent tools**, paired with an external **Counterfactual Audit Protocol** that grades agent submissions without using the library to grade itself. Across synthetic worlds with 12 planted defects (honest ceiling Sharpe 1.63 vs. same-day shortcut 71.5), our guards achieve **100% recall** (36/36) with **zero false alarms**, while physical session-by-session clock replay exposes leaky pipelines that score Sharpe 22.80 on held-out static years yet lose -1.17 bps/day (Sharpe -0.58) live. In a multi-model agent evaluation, **fin-skills** reduces strong-model (Claude Opus) Sharpe reporting error by **84%** ($0.19 \rightarrow 0.03$), but uncovers a striking **Compliance Paradox**: giving weak models (Claude Haiku) passive skill documentation alone yields the worst distortion (11.95 Sharpe gap) as agents *hallucinate citations to guards they never executed*. Finally, auditing **1.72M bilingual posts and 1,445 A-share KOL track records** reveals a $3.96\times$ follower asymmetry that drives unweighted NLP into negative correlation (Rank IC = -0.0318 , $p = 1.74 \times 10^{-20}$) until reversed by Bayesian credibility gating (Rank IC = $+0.0104$, $p = 0.001$) and LLM semantic distillation ($+0.0084 \rightarrow +0.0186$).

1 Introduction

Equipped with code interpreters and financial data feeds, LLM agents can formulate hypotheses, engineer alpha signals, and report backtest performance in minutes (Wu et al., 2023; Yang et al., 2023; Yu et al., 2024). Yet in quantitative finance, high backtest Sharpe ratios almost never survive live capital deployment (Bailey and López de Prado, 2014; Bailey et al., 2017; López de Prado, 2018). Through systematic auditing of LLM research trajectories, we show that agentic backtest inflation stems from two orthogonal failure classes:

- 1. Type I: Mechanical Research Shortcuts.** Agents silently exploit structural artifacts in tabular workspaces: joining post-close news commentary to the same session’s open-to-close return, training on pre-computed LLM sentiment scores whose fitting window overlaps the evaluation period, filtering universes by future survival, or computing factors on unadjusted split prices.
- 2. Type II: Empirical & Microstructure Illusions.** Even when timestamps are strictly causal, agents misinterpret noisy social streams and market rules. On financial social platforms, viral retail sentiment peaks at local price tops, while real execution is governed by $T + 1$ settlement lockups, 100-share board lots, asymmetric sell-side stamp duties, and secondary-market ETF premium caps.

We present **fin-skills** paired with an external **Counterfactual & Microstructure Audit Protocol**. Rather than claiming to manufacture market alpha, our framework enforces verifiable research integrity:

- **The fin-skills Verification Library (§3):** 129 source-verified skills, 36 executable guards (`Bundle.check()`), and 53 Model




Figure 1: **Overview of our two core empirical discoveries.** (a) **The Agent Compliance Paradox (§5):** Equipping a strong model (Claude Opus) with *fin-skills* reduces Sharpe reporting error by 84% (0.19 $\rightarrow$ 0.03), whereas giving a weak model (Claude Haiku) passive text documentation without mandatory runtime execution induces severe reporting distortion ($|\Delta \text{Sharpe}| = 11.95$) due to *hallucinated guard citations*. (b) **The $3.96 \times$ KOL Megaphone Bias (§5.2):** Across 1,445 verified A-share influencers, statistically confirmed contrarian indicators ($N = 98$, 5-day Rank IC -0.2130) command $3.96 \times$ more followers (55.99k) than Tier-1 Alpha researchers ($N = 70$, 14.12k fans, Rank IC $+0.2158$), explaining why unweighted financial NLP loses money.

Context Protocol (MCP) tools validated by 3,086 unit tests, achieving 8/8 numerical parity against `scipy/statsmodels` and $\mathcal{O}(N^{0.2637})$ sub-linear runtime scaling.

- **External Counterfactual Auditing (§4):** Four black-box perturbation operators and session-by-session clock replay that detect look-ahead and same-session leakage without inspecting agent code or using the library to grade itself.
- **The Compliance Paradox (§5):** Controlled 3-condition agent ablations proving that passive prompt guidelines induce weak models to hallucinate guard execution, whereas runtime provenance gating eliminates fabricated citations (2/2 $\rightarrow$ 0/3).
- **Real-World Bilingual KOL & Microstructure Audit (§5.2):** A self-contained reproducibility bundle over 1.72M bilingual posts and 1,445 A-share KOL track records that quantifies the $3.96 \times$ follower asymmetry paradox and reverses negative NLP correlation (Rank IC : $-0.0318 \rightarrow +0.0104$).

2 Related Work

Financial LLMs & Autonomous Agents. Recent domain foundation models (Wu et al., 2023; Yang et al., 2023) and memory-augmented trading agents (Yu et al., 2024) demonstrate strong capabilities in financial sentiment extraction and strategy synthesis. However, existing evaluations rely primarily on static historical splits that remain vulner-

able to pre-training cutoff leakage, same-session timestamp alignment errors, and retail social echo-chambers.

Backtest Overfitting & Behavioral Verification.

Econometric controls for selection bias—including White’s Reality Check (White, 2000), Hansen’s SPA test (Hansen, 2005), and the Deflated Sharpe Ratio (Bailey and López de Prado, 2014; Bailey et al., 2017)—quantify multiple-testing inflation after honest returns are computed, but cannot detect mechanical code bugs or future-leaking table joins. Inspired by behavioral testing in NLP (Ribeiro et al., 2020) and modular verification toolkits (Zhao et al., 2019), *fin-skills* pairs machine-verifiable domain specifications with black-box counterfactual perturbation oracles.

3 The *fin-skills* Verification Architecture

fin-skills unifies domain specifications and executable code verification across seven packages (`fin-core`, `fin-ml`, `fin-market-data`, `fin-alt-data`, `fin-china`, `fin-strategies`, and `fin-llm`):

Source-Verified Skills (129 Specifications).

Each skill obeys a machine-validated 6-field schema carrying explicit verification timestamps (`verified_on`) and primary regulatory or academic citations. Skills pinpoint silent traps in standard quantitative libraries, including fractional differentiation truncation (López de Prado, 2018), day-count conventions, and point-in-time SEC Form 4 /

13F acceptance timestamps.

Executable Guards (`Bundle.check()`) & MCP Tools. To bridge passive reading and run-time enforcement, 36 executable Python guards operate over a unified slot-binding contract:

$$\mathcal{R} = \text{Bundle}(\mathbf{r}, \mathbf{P}, \tau, \mathcal{C}).\text{check}() \in \{\text{PASS}, \text{WARN}, \text{FAIL}\} \quad (1)$$

where $\mathbf{r}$ denotes strategy returns, $\mathbf{P}$ OHLCV price bars, τ turnover, and $\mathcal{C}$ point-in-time corporate action and filing timestamps. All 36 guards and 8 econometric/portfolio primitives export native MCP schemas (53 tools total). In Experiment E3 (`benchmarks/parity_and_cost_bench.py`), our primitives achieve **8/8 exact numerical parity** against `scipy` and `statsmodels` (maximum error 1.14×10^{-8} on SLSQP minimum-variance QP, 6.94×10^{-18} on analytical Expected Shortfall, and 2.12×10^{-10} on Augmented Dickey-Fuller statistics). Across $N \in \{250, 1000, 5000, 25000\}$ rows, `Bundle.check()` scales sub-linearly as $\mathcal{O}(N^{0.2637})$, executing all applicable guards in **111.17 ms P50** (3.66 MB peak RAM) at $N = 25,000$ rows (~ 100 trading years).

4 External Counterfactual Audit Methodology

A core methodological principle of our study is that *an evaluation harness must never grade agent submissions using the same library guards given to the agent*. We construct an independent ground-truth **Seeded Market Generator** (`benchmarks/agent_study/`) spanning 60 equities over 2017–2023 with four planted shortcuts: (i) a news feed stamped with the session it describes, (ii) an LLM sentiment score trained through 2020-12-31, (iii) a listing table exposing future delisting dates, and (iv) unadjusted split prices.

Honest Ceiling vs. Shortcut Inflation. Because the data-generating process is known in closed form, we compute the exact theoretical ceiling: trading the generator’s true latent drift with a strict 1-day lag earns **Net Sharpe 1.63**. By contrast, joining same-session news earns **Net Sharpe 71.5**. Across three multi-seed synthetic worlds ($s \in \{17, 29, 43\}$, 36 planted defects total), our guards achieve **100% recall** (36/36) with **0 false alarms** across 400 clean control runs (where planted flaws inflate reported Sharpe from -0.052 ± 0.309 to $+2.210 \pm 0.562$).

Reference Pipeline	Leakage	Paper OOS	Live Clock Rep
Honest (lagged feed, split-adj.)	0.00	1.63	+5.07 bps/d (SR)
Leaky (same-day join, full scaler)	0.78	22.80*	−1.17 bps/d (SR)

Table 1: Instrument calibration (*illusory paper Sharpe). Static year holdout (2023) fails to punish same-session leakage (22.80); physical clock replay exposes it

Agent Arm ($n = 3/\text{cell}$)	$ \Delta\text{Sharpe} $	Leakage	Same-Day	Contam. I
Claude Opus (No Library)	0.19	0.00	0.00	0/3
Claude Opus + <code>fin-skills</code>	0.03	0.00	0.00	0/3
Claude Haiku (No Library)	0.57	0.00	0.00	2/3
Claude Haiku + <code>fin-skills</code>	11.95	0.00	0.42	2/3

Table 2: Multi-model pilot across seeds $\{11, 23, 37\}$. $|\Delta\text{Sharpe}| = |\text{Claimed} - \text{Recomputed}|$. Weak model + passive library text yields severe distortion (11.95, Fig. 1a).

Four Black-Box Counterfactual Perturbations.

Given any agent-submitted weight function $\mathbf{W} = f(\mathcal{D})$, our external oracle (`oracle.py`) measures four invariants without reading agent source code:

- Future Redraw (leakage_rate):** Redraw market paths after cut date t_c ; compute the fraction of pre-cut weights $\mathbf{W}_{t \leq t_c}$ that change across 3 redraws.
- Single-Session Redraw (same_session_rate):** Redraw session t alone and inspect $\mathbf{W}_t$, catching same-day post-close joins invisible to cut-date holdouts.
- Input Scramble (uses_(file)):** Permute individual input tables to test causal reliance on contaminated features (`llm_score.csv`).
- Physical Clock Replay:** Withhold future sessions physically by feeding data day-by-day (`live_replay.py`).

5 Experiments

5.1 Agent Evaluation & The Compliance Paradox

Table 1 proves that holding out a static future year (2023) does not prevent same-session leakage (22.80 paper Sharpe vs. -0.58 live clock replay). Table 2 and Fig. 1a reveal three core findings: **(1) Strong Models Gain Reporting Precision:** Claude Opus avoided all four traps independently, and `fin-skills` reduced its Sharpe reporting error by **84%** ($0.19 \rightarrow 0.03$). **(2)**

KOL Credibility Tier	Count	Mean Fans	5D Win%
TIER_1_CORE_ALPHA	70	14,121.9	68.47%
TIER_2_RESEARCHER	115	31,994.4	70.92%
TIER_0_ELITE_KOL	31	186,610.1	61.65%
TIER_MEDIA_AGGREG.	11	107,354.3	46.98%
TIER_NEUTRAL_RETAIL	1,120	71,251.4	46.72%
TIER_CONTRARIAN	98	55,988.6	28.49%

Table 3: Stratification of 1,445 A-share KOLs. Contrarian indicators average $3.96\times$ more followers than Tier-1 Alpha researchers (55,989 vs. 14,122, Fig. 1b).

Model / Representation Variant	AUC	Acc.
MMAN Unweighted Fact Surge	0.4850	48.21%
MMAN Baseline (Price + Raw Text)	0.5016	49.82%
MMAN + Bayesian Credibility Gate	0.5033	51.13%
Chinese RoBERTa (Raw Social Text)	0.5075	50.67%
Chinese RoBERTa (LLM-Distilled)	0.5121	50.42%

Table 4: Out-of-sample predictive ablation ($N = 334,946$ MMAN; $N = 31,505$ RoBERTa). Credibility gating reverses negative Rank IC ($\Delta\text{IC} = +0.04224$); LLM distillation doubles RoBERTa IC ($p < 0.01$).

The Compliance Paradox: Equipping Claude Haiku with skill documentation produced the **worst reporting distortion** (11.95 mean $|\Delta\text{Sharpe}|$), claiming Sharpe 0.24 while earning 27.96 in one run and claiming 0.13 while losing -4.86 in another. Crucially, our 3-condition scaffold (`open_agent_runner.py`) proves that under text-only guidance (Condition B), weak agents *hallucinate citations to guards* (`safe_asof`, `contamination_probe`) in their final reports without ever executing them (2/2 fabricated citations caught), whereas mandatory runtime provenance gating (Condition C) eliminates hallucinated compliance (0/3). **(3) Integrity Tools Do Not Manufacture Alpha:** Across held-out 2023 live replays, library access yielded 1 win ($+4.44\%$ vs. $+3.90\%$), 1 loss (-0.64% vs. $+2.36\%$), and 1 tie ($+0.30\%$ vs. $+0.55\%$), confirming that guards prevent false discoveries rather than inventing edge.

5.2 Real-World Bilingual KOL & Microstructure Audit

To evaluate Type II empirical and microstructure illusions in production markets, we audit **1.72M bilingual social posts** (Chinese Xueqiu and US StockTwits) paired with **1,445 verified A-share KOL track records** (≥ 15 forward-priced calls each; `benchmarks/real_world_kol_audit.py`).

The Rank IC × Follower Asymmetry & Bayesian Reversal Table 3 and Fig. 1b explain why standard financial NLP fails on social streams: accounts in TIER_CONTRARIAN_INDICATOR ($N = 98$, mean 5-day IC = -0.2130 , win rate 28.49%) command **55,988.6 average followers**— $3.96\times$ more than **TIER_1_CORE_ALPHA** researchers (14,121.9 followers, IC = $+0.2158$). Consequently, unweighted Multimodal Attention Networks (MMAN) suffer severe negative correlation across 334,946 test samples (Rank IC = -0.03184 , $p = 1.74 \times 10^{-20}$, Table 4). Enforcing `fin-kol-credibility-registry` skill—`onjugate` Beta-Binomial shrinkage with polarity inversion for statistically verified contrarians—**reverses** predictive performance to Rank IC = $+0.01040$ ($p = 0.001$, $\Delta\text{IC} = +0.04224$). Moreover, stripping emotional hype via LLM semantic distillation ($N = 31,505$, reducing the hype-to-hedging ratio by 79.8% and boosting quantitative fact density by 218%) more than doubles standalone Chinese RoBERTa Rank IC from $+0.00842$ ($t = 1.49$) to $+0.01862$ ($t = 3.30$).

A-Share Microstructure Execution Gates. Deploying six `fin-china` execution guards ($T + 1$ lockup, 100-share board-lot rounding, asymmetric sell-side stamp duty, QDII ETF premium cap $\leq 1.5\%$, cash drag, and Tier-C noise-stock blocking) in an 80/20 Core-Satellite portfolio (2023–2026 OOS) improves annualized Sharpe to **2.76** (**1.34** conservative baseline) and slashes maximum drawdown from -27.65% (CSI 300) and -8.85% (unfiltered satellite) to -3.94% (Calmar **5.74**).

6 Conclusion

We introduced **fin-skills** alongside an external counterfactual and microstructure audit methodology for financial AI agents. By combining black-box counterfactual perturbations, runtime guard provenance tracking, and Bayesian KOL credibility calibration, our framework eliminates backtest shortcut illusions and bridges the gap between financial NLP benchmarks and live market execution.

Limitations

In accordance with rigorous task-method scope governance, we explicitly bound the operating regime and failure boundaries of `fin-skills`:

1. **Integrity Verification vs. Alpha Generation:** As demonstrated in §5.1, `fin-skills` is an integrity verification toolkit rather than an alpha-generation engine. In held-out live replays where underlying strategies lack predictive edge, guard enforcement prevents false discoveries without artificially increasing realized returns.
2. **Sample Size of Open-Weight Agent Sweep:** While our controlled multi-model evaluation and 3-condition verification scaffold establish clear discrimination between strong and weak agents and expose hallucinated guard citations, scaling statistical power across all open-weight model families requires completing our pre-registered full factorial sweep (Experiment E2: 12 markets $\times$ 3 model sizes $\times$ 3 conditions $\times$ 3 repetitions, detailed in Appendix B).
3. **Non-Stationarity of Social Platform Algorithms:** Real-world social sentiment dynamics evolve as platforms modify recommendation algorithms, requiring periodic rolling re-estimation of Bayesian KOL priors to avoid structural drift.

Ethics Statement

All 1,445 social media author profiles in our released benchmark fixture (`benchmarks/data/kol_credibility_linguistic_fixture.json`) are cryptographically pseudonymized via one-way SHA-256 digests (`KOL_<10-hex>`) with raw author handles, URLs, and personally identifiable attributes permanently stripped to protect individual privacy. None of the skills, guards, or empirical analyses in this work constitute investment advice or automated live-money order execution; all broker-facing skills enforce strict paper-account sandboxing (`paper_account_guard`) and human-in-the-loop pre-trade kill switches. By exposing fabricated backtest Sharpe ratios and hallucinated audit compliance, our work directly protects retail and institutional participants from deploying leaky AI-generated trading strategies.

References

- David H. Bailey, Jonathan Borwein, Marcos López de Prado, and Qiji Jim Zhu. 2017. The probability of backtest overfitting. *Journal of Computational Finance*, 20(4):39–69.
- David H. Bailey and Marcos López de Prado. 2014. The deflated Sharpe ratio: Correcting for selection bias, backtest overfitting, and non-normality. *Journal of Portfolio Management*, 40(5):94–107.
- Peter Reinhard Hansen. 2005. A test for superior predictive ability. *Journal of Business & Economic Statistics*, 23(4):365–380.
- Marcos López de Prado. 2018. *Advances in Financial Machine Learning*. John Wiley & Sons.
- Marco Tulio Ribeiro, Tongshuang Wu, Carlos Guestrin, and Sameer Singh. 2020. Beyond accuracy: Behavioral testing of NLP models with CheckList. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 4902–4912.
- Halbert White. 2000. A reality check for data snooping. *Econometrica*, 68(5):1097–1126.
- Shijie Wu, Ozan Irsoy, Steven Lu, Vadim Dabravolski, Mark Dredze, Sebastian Gehrmann, Prabhanjan Kam-badur, David Rosenberg, and Gideon Mann. 2023. BloombergGPT: A large language model for finance. *arXiv preprint arXiv:2303.17564*.
- Hongyang Yang, Xiao-Yang Liu, and Christina Dan Wang. 2023. FinGPT: Open-source financial large language models. *FinLLM Symposium at IJCAI*.
- Yangyang Yu, Haohang Li, Zhi Chen, Yuechen Jiang, Yang Li, Denghui Zhang, Rong Liu, Jordan W. Su-chow, and Khaldoun Khashanah. 2024. FinMem: A performance-enhanced LLM trading agent with layered memory and character design. In *Proceedings of the AAAI Symposium Series*, volume 3, pages 595–597.
- Yue Zhao, Zain Nasrullah, and Zheng Li. 2019. PyOD: A Python toolbox for scalable outlier detection. *Journal of Machine Learning Research*, 20(96):1–7.

A Experiment E1: Synthetic Defect Recall Across Random Seeds

In `benchmarks/guard_bench.py`, we evaluate guard recall across three independent synthetic market worlds generated with random seeds $s \in \{17, 29, 43\}$, each containing 12 planted research defects spanning look-ahead joins, unadjusted corporate splits, survivorship filtering, non-monotonic as-of merges, cost implausibility, and regime look-ahead. Across all three worlds, **fin-skills** achieved **100% recall** (36/36 total defects detected), **0 false alarms** across 400 clean control runs, and **0 runtime exceptions** (`benchmarks/GUARD_ROBUSTNESS.json`).

B Experiment E2: Three-Condition Open-Weight Scaffold & Compliance Auditor

To strictly isolate whether agent improvements stem from passive prompt documentation or active runtime execution, we implement `benchmarks/agent_study/open_agent_runner.py`, which evaluates any OpenAI-compatible open-weight server (Qwen2.5-Coder-32B-Instruct, DeepSeek-V3, Llama-3.3-70B-Instruct) across three isolated conditions:

- **Condition A (no_library):** Task prompt and raw market CSVs only.
- **Condition B (skills_text_only):** Injects relevant `SKILL.md` guidance into context while physically blocking `fin_skills.api` guard imports.
- **Condition C (skills_plus_guards):** Provides both `SKILL.md` text and executable guard tools with immutable runtime provenance logs (`guard_execution_trace.json`).

Our ComplianceAuditor cross-checks every guard name cited in the agent’s final report against the runtime provenance log, flagging any cited guard that was not actually invoked on the submission outputs (`SCAFFOLD_VERIFICATION_REPORT.json`).

C Experiment E3: Complete Numerical Parity & Guard Scaling Tables

Table 5 reports the exact numerical parity verification across all 8 financial and econometric

primitives, alongside the full wall-clock timing and peak RAM profile of `Bundle.check()` across dataset scales (`benchmarks/PARITY_AND_COST_RESULTS.json`).

Primitive	Reference Benchmark	Max $ \Delta $	Status
Equal Weight	Exact closed form ($1/N$)	0.00	PASS
Inverse Vol.	Exact sample σ_i^{-1}	0.00	PASS
Min-Var QP	scipy.SLSQP convex QP	1.14×10^{-8}	PASS
Normal VaR/ES	scipy.stats.norm exact	6.94×10^{-18}	PASS
Hist. VaR/ES	Empirical quantile & tail mean	0.00	PASS
Hist. Vol.	pandas sample std ($ddof = 1$)	0.00	PASS
Frac. Diff. FFD	Closed-form binomial ($d = 0.4$)	1.11×10^{-16}	PASS
ADF Unit Root	statsmodels.adfuller	2.12×10^{-10}	PASS
Guard Suite Runtime Scaling ($\mathcal{O}(N^{0.2637})$ across 36 guards)			
$N = 250$ rows	P50: 32.27 ms (P95: 36.02 ms)	RAM: 0.14 MB	PASS
$N = 1,000$ rows	P50: 44.08 ms (P95: 47.91 ms)	RAM: 0.24 MB	PASS
$N = 5,000$ rows	P50: 90.40 ms (P95: 96.85 ms)	RAM: 0.79 MB	PASS
$N = 25,000$ rows	P50: 111.17 ms (P95: 118.42 ms)	RAM: 3.66 MB	PASS

Table 5: Complete E3 numerical parity (8/8 passed) and runtime scaling profile.

D Experiment E4: Linguistic Feature Shifts & Standalone Reproducibility

Table 6 summarizes the pragmatic and lexical transformations achieved by LLM semantic distillation over $N = 31,505$ noisy financial social posts (`benchmarks/real_world_kol_audit.py`). Every metric in §5.2 can be reproduced in $< 1s$ on a single CPU core via `real_world_kol_audit.py`.

Linguistic Dimension	Raw Posts	Distilled	Shift
Quant. Fact Density (/100c)	1.049	1.078	+2.8% (Tok +218%)
Hype-to-Hedging Ratio	4.85	0.98	−79.8%
Sarcasm / Polarity Error	18.4%	3.1%	−83.2%
Ticker Attribution Prec.	61.4%	96.2%	+34.8 pp

Table 6: Linguistic feature transformation explaining the $2.21\times$ Rank IC improvement in Chinese RoBERTa.




Figure 2: **System scaling and counterfactual clock replay calibration.** (a) Across dataset sizes $N \in [250, 25000]$ rows, `Bundle.check()` scales sub-linearly as $\mathcal{O}(N^{0.2637})$, completing in 111.17 ms P50 with 3.66 MB peak memory at $N = 25,000$ rows. (b) Static future-year holdout (2023) fails to punish same-session post-close joins (Sharpe = 22.80), whereas physical session-by-session clock replay (`live_replay.py`) exposes the true live loss (Sharpe = -0.58).